

Task 3. Reinforcement Learning

Garoe Dorta-Perez
CM50245: Computer Animation and Games II

March 17, 2015

1 Introduction

2 A* path finding

The A* algorithm can find paths optimally in a graph based on a heuristic. It works under the assumption that the distance heuristic will never overestimate the cost of the path and the path cost will not decrease as we travel through it. In a maze world, where the movement possibilities are north, south, east and west, with fixed positive costs. Using a Manhattan distance for the actual path cost and a euclidean distance to compute the heuristic will fit the assumptions.

In Algorithm 1 an overview of a general implementation of the A* in pseudocode is given. As stated in the previous paragraph, in our case, `graph.cost()` and `heuristic()` returns respectively, the Manhattan and the euclidean distance between two nodes.

3 Q learning in noughts and crosses

4 Q learning in BlackJack

5 Results and conclusions

Algorithm 1: A*

Data: *goal* goal position, *start* start position, *graph* graph with the tiles in the map.

Result: *path* path from *start* to *goal*, *cost* total cost of the path.

frontier = PriorityQueue();

frontier.put(start, 0);

came_from = {};

cost_so_far = {};

came_from[start] = None;

cost_so_far[start] = 0;

while not frontier.empty() **do**

 current = frontier.get();

if current == goal **then**

 | break;

end

for next in graph.neighbors(current) **do**

 new_cost = cost_so_far[current] + graph.cost(current, next);

if next not in cost_so_far or new_cost < cost_so_far[next] **then**

 cost_so_far[next] = new_cost;

 priority = new_cost + heuristic(goal, next);

 frontier.put(next, priority);

 came_from[next] = current;

end

end

end

path = getPath(came_from);

goal = cost_so_far[current];
